

Gated scores — eight frontier models, three vendors, six gold cases

▲ = auto-fail gate fired (GATE.<name>); one run per model per case

Claude Opus 4.8	0.965	0.974	0.983	0.983	0.980	0.980	flagship
Claude Sonnet 4.6	0.955	0.973	0.943	0.983	0.933	0.980	mid
Claude Haiku 4.5	0.692 ▲ C1FCF	0.799	0.496 ▲ SCALE	0.983	0.933	0.980	small
GPT-5.6-sol	0.951	0.970	0.943	0.973	0.980	0.980	flagship
GPT-5.5	0.953	0.970	0.943	0.983	0.980	0.980	flagship (prev.)
GPT-5.4	0.903	0.883	0.983	0.973	0.980	0.980	mid
GPT-5.4-mini	0.631 ▲ FALSEPREC.	0.650 ▲ BRIDGE	0.714	0.983	0.933	0.980	small
Gemini 3.6 Flash	0.922	0.970	0.983	0.973	0.980	0.980	small/fast
	#3 DCF (MCD)	#3 DCF (NVDA)	#4 recon (break)	#4 recon (clean)	#5 confirm (break)	#5 confirm (clean)	